

Automated Multi-Object Clustering Using Machine Learning

Learning to select clustering algorithms from geometric and statistical dataset descriptors

Authors: Hagen Jacob, Jan Edgar König, Robert Vetter **Affiliation:** Spezialschulenteil des Albert-Schweitzer-Gymnasiums Erfurt (specialised STEM track) **External scientific guidance:** Dr. Wolfgang Felber and Christopher Sobel, Fraunhofer IIS, Nuremberg **Type:** Secondary-school research project (*Seminarfacharbeit*, grades 11/12, 2022–2024)

This is a student/seminar research report. It is **not** peer-reviewed and is not a journal or conference publication. It is an English write-up of the original German seminar paper ([seminar-paper.pdf](#)), with results reproduced from the project's code and committed data. All reported metrics are on **synthetic** data unless stated otherwise.

Abstract

Choosing an appropriate clustering algorithm for an unlabelled dataset is normally a manual, expertise-dependent task: different algorithms suit different geometries, and there is no universally best choice. This project investigates whether the choice can be **automated** by describing a dataset with a small set of geometric and statistical descriptors and learning a mapping from those descriptors to the most suitable algorithm. We build a three-stage pipeline — density-based outlier removal (DBSCAN + Local Outlier Factor), contour-based shape characterisation (Gaussian kernel-density estimation and a histogram-edge method, selected by silhouette score), and extraction of **ten** geometric/ statistical features — followed by a compact neural-network classifier ($\approx 2,500$ parameters) that selects among four complementary algorithms: K-Means, DBSCAN, OPTICS, and Gaussian Mixture Models. The classifier is trained on **5,184 synthetic datasets**, balanced across the four classes. On a held-out synthetic test split it reaches **99.81% accuracy** with near-uniform per-class F1-scores (≈ 0.998). We report the methodology, experiments, and results in detail, and are explicit about the central limitation: the evaluation is on synthetic data, where each dataset's label follows from its generating family; generalisation to arbitrary real-world data is markedly lower (the companion article reports $\approx 70\%$) and is left to future work.

Keywords: clustering, algorithm selection, meta-features, neural networks, outlier detection, unsupervised learning.

1. Introduction and motivation

Clustering partitions unlabelled data into groups of similar points and is a core tool of unsupervised learning. In practice, however, the analyst must first decide *which* clustering algorithm to use: K-Means assumes compact, roughly spherical clusters; DBSCAN finds arbitrarily-shaped, density-connected clusters and is robust to noise; OPTICS extends density-based clustering to varying densities; and Gaussian Mixture Models (GMM) model soft, overlapping, elliptical clusters. A poor choice produces misleading groupings.

Selecting the algorithm is itself a decision problem that is usually solved by hand using domain expertise and trial-and-error. This project asks whether that decision can be **learned**: given a dataset, can a model recommend the best-suited clustering algorithm automatically? Concretely, we restrict the

choice to four complementary algorithms and learn a classifier from a compact, fixed-length description of a dataset's geometry and statistics.

This is an instance of the **algorithm-selection problem** [Rice, 1976]: map a problem instance, via measurable features, to the best-performing algorithm from a portfolio. Our contribution is a concrete, end-to-end pipeline for the clustering case, in which the dataset-description step is itself non-trivial and combines density estimation, contour extraction, and classical shape statistics.

2. Background and related work

The algorithm-selection problem. Rice [1976] framed algorithm selection as learning a mapping from a feature space describing problem instances to a performance space over a portfolio of algorithms. Modern meta-learning and AutoML follow this template, characterising datasets by *meta-features* and learning which method works best. This project applies that idea to clustering, with hand-designed geometric/statistical meta-features.

Clustering algorithms. The four algorithms in our portfolio are: K-Means [Lloyd, 1982; MacQueen, 1967], DBSCAN [Ester et al., 1996], OPTICS [Ankerst et al., 1999], and Gaussian Mixture Models fitted via Expectation–Maximisation [Dempster et al., 1977]. They are complementary: centroid-, density-, ordering-, and probability-based respectively.

Outlier detection. Outliers distort geometric descriptors, so they are removed first. We combine DBSCAN's intrinsic noise labelling with the Local Outlier Factor (LOF) [Breunig et al., 2000], a local density-deviation score.

Cluster validity and density estimation. Internal choices (DBSCAN `min_samples`, the density threshold, the smoothing parameter) are made with the silhouette coefficient [Rousseeuw, 1987]. Shape is captured from a kernel density estimate [Rosenblatt, 1956; Parzen, 1962].

3. Problem formulation

Let a dataset be a finite point set $X = x_1, \dots, x_n \subset \mathbb{R}^2$. Let $\mathcal{A} = \text{KMeans}, \text{DBSCAN}, \text{OPTICS}, \text{GMM}$ be the algorithm portfolio. We seek a function

$$F : X \mapsto a \in \mathcal{A}$$

that returns the algorithm best suited to X . We factor F as


$$F = c \circ \phi, \quad \phi : X \rightarrow \mathbb{R}^{10}, \quad c : \mathbb{R}^{10} \rightarrow \mathcal{A},$$

where ϕ maps a dataset to a fixed 10-dimensional feature vector (Section 4.4) and c is a neural-network classifier (Section 4.5). The map ϕ is the core difficulty: it must summarise an arbitrary-size point cloud into a small, informative, fixed-length descriptor.

4. Methodology

The full pipeline is shown in Figure 1: outlier removal $\rightarrow$ shape recognition $\rightarrow$ feature extraction $\rightarrow$ classification.

Figure 1. *The four-stage pipeline.*



4.1 Synthetic dataset generation

Training data is generated synthetically with scikit-learn's dataset generators. Four generators are used, each associated with the clustering algorithm best suited to its structure **by construction**:

Generator (<code>sklearn.datasets</code>)	Structure	Assigned label
<code>make_blobs</code>	compact, round clusters	KMeans
<code>make_circles</code>	concentric rings	OPTICS
<code>make_moons</code>	interleaving crescents	DBSCAN
<code>make_gaussian_quantiles</code>	nested Gaussian shells	GMM

Datasets are produced over grids of parameters (sample counts in $\{500, 800, 1000, 1200, 1500, 1700, 2000, 2500, 4000\}$; cluster counts $\{3, 4, 5, 6\}$; noise levels $\{0.02, 0.04, 0.06\}$; circle factors $\{0.3, 0.5\}$; quantile classes $\{1, 2, 3\}$). Each base dataset is additionally **augmented** (Section 4.1.1) to increase variety and robustness. The procedure yields **5,184 datasets**, exactly **balanced** at **1,296 per class**.

For each dataset the 10-feature vector (Section 4.4) is computed; the resulting table is scaled **column-wise** with a `RobustScaler` (median/IQR, robust to outliers) and stored as the training table `data/data_scaled.xlsx`.

Note on labels. The label is the algorithm assumed best for the generating family, not the empirical winner of running all four algorithms on each dataset. This is a deliberate simplification and a key limitation (Section 8).

4.1.1 Data augmentation

Each base point cloud is transformed by a randomised, label-preserving sequence (class `DataAugmentor`): removing a random subset of points; adding jittered copies of existing points ($+0.05\mathcal{N}(\mathbf{0}, \mathbf{I})$); isotropic scaling by a factor in $[0.5, 1.5]$; small uniform shifts; and a random rotation about the centroid by $\theta \sim \mathcal{U}[0, 2\pi]$.

4.2 Anomaly detection (Stage 1)

Outliers are removed so they do not distort the shape descriptors. Two detectors are combined and their flagged points unioned.

DBSCAN with an automatic radius. For a candidate `min_samples =  $k$` , the radius is set by the k -distance heuristic: let $d_k(x)$ be the distance from x to its k -th nearest neighbour; then

$$\varepsilon = Q_{0.93}(\{d_k(x) : x \in X\}),$$

the 93rd percentile, which admits at most $\sim 7\%$ of points as low-density (noise). `min_samples` is swept over $[2d, 3d]$ for data dimension d (here $d = 2$), and the value maximising the silhouette score of the resulting non-noise clusters is kept.

Local Outlier Factor. LOF [Breunig et al., 2000] scores each point by how isolated it is relative to its neighbours. Its hyperparameters (`n_neighbors` $\in [10, 25]$, `contamination` $\in \{0.01, 0.02, 0.03\}$) are grid-searched against the known synthetic inlier/outlier labels using the **macro F1-score**, appropriate for the strong class imbalance between inliers and outliers.

The final outlier set is the union of the DBSCAN-noise points and the LOF-flagged points; these are removed before Stage 2.

4.3 Shape / contour recognition (Stage 2)

The geometry of the cleaned cloud is captured by the contours enclosing its dense regions. Two independent methods are computed and the better one is kept by silhouette score.

Method 1 — density contours. A 2-D Gaussian kernel-density estimate [Parzen, 1962] is evaluated on a 100×100 grid,

$$\hat{f}(z) = \frac{1}{n} \sum_{i=1}^n K_H(z - x_i),$$

and contour lines are drawn at the level $\lambda = \text{factor} \cdot \hat{f}$ for several factors $\in \{0.7, 1.0, 1.5, 3.0\}$; the factor maximising the silhouette score (Section 4.3.1) is chosen.

Method 2 — histogram edges. A padded 64×64 2-D histogram of the points is convolved with a Sobel operator to obtain an edge-magnitude field $\|\nabla h\|$, smoothed by a Gaussian filter of width σ , thresholded, and run through `skimage.measure.find_contours`. The contours are rescaled to data coordinates. The smoothing $\sigma \in \{0.7, 1, 1.5, 3\}$ that maximises the silhouette score is chosen.

Method selection. The method whose contours give the higher silhouette score is used downstream (Figures 2–3).

4.3.1 Contour scoring

To score a candidate set of contours, each contour is treated as a polygon and every point is assigned the index of the polygon that contains it (or marked as noise if contained by none). The silhouette coefficient [Rousseeuw, 1987] of that labelling,

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}},$$

(with $a(i)$ the mean intra-cluster distance and $b(i)$ the mean nearest-cluster distance), is averaged over points and used as the contour-quality score.

Figure 2. An example synthetic dataset (left) and the extracted contours overlaid on it (right).





Figure 3. Filled shape recognition: density method (left) vs. histogram method (right). The better-scoring method is selected per dataset.



4.4 The ten features (Stage 3)

The cleaned cloud and its selected contours are reduced to a fixed 10-dimensional vector. The contours give shape descriptors; the points give distributional descriptors.

#	Feature	Definition
1	mean eccentricity	mean over contours of $e = \sqrt{1 - (b/a)^2}$, the eccentricity of the best-fit ellipse (axes $a \geq b$) via <code>cv2.fitEllipse</code>
2	std eccentricity	standard deviation of the per-contour eccentricities
3	contour/convex-hull area ratio	mean over contours of (contour area)/(convex-hull area); $\rightarrow 1$ for convex shapes
4	variance of centroid distances	$\text{Var}(\ x_i - \bar{x}\)$ over points
5	contour-area/points ratio	mean over contours of (polygon area)/(number of enclosed points)
6	mean kurtosis	mean of the per-axis kurtosis of the coordinates
7	std kurtosis	standard deviation of the per-axis kurtosis
8	mean skewness	mean of the per-axis skewness of the coordinates
9	std skewness	standard deviation of the per-axis skewness
10	compactness	$\frac{1}{n} \sum_i \ x_i - \bar{x}\ ^2$, mean squared distance to the centroid

4.5 Neural-network classifier (Stage 4)

The classifier $\mathcal{C}$ is a compact Keras `Sequential` network:

```
Input(10)
Dense(48, ReLU, L2=0.01) → Dropout(0.5) → BatchNorm
Dense(32, ReLU, L2=0.01) → Dropout(0.5) → BatchNorm
Dense(4, softmax)
```

Class probabilities use the softmax $\sigma(z)_j = e^{z_j} / \sum_k e^{z_k}$, and training minimises the (sparse) categorical cross-entropy $\mathcal{L} = -\sum_j y_j \log \hat{y}_j$. Optimisation uses Adam with learning rate 10^{-4} , batch size 32, for 100 epochs. The model has **2,548 parameters** (2,388 trainable), of which the dropout (0.5), L2 weight decay, and batch normalisation provide regularisation against overfitting on the small feature table.

5. Experimental setup

Item	Value
Dataset	5,184 synthetic datasets → 10 features each, <code>RobustScaler</code> -scaled
Class balance	1,296 each: KMeans / DBSCAN / OPTICS / GMM
Split	80% train / 20% test, <code>train_test_split(random_state=42)</code>
Validation	20% of the training set (<code>validation_split=0.2</code>)
Resulting sizes	4,147 train (incl. validation) / 1,037 test
Optimiser	Adam, lr 10^{-4} , batch size 32, 100 epochs
Seed	42 (<code>tf.keras.utils.set_random_seed</code>) for the run reported here
Environment	Python 3.11, TensorFlow 2.x, scikit-learn (see Appendix A)

The results below come from a single reproducible run (seed 42) using the committed data and the architecture above. For reference, the original project notebook independently logged $\approx 99.90\%$ test accuracy on its (unseeded) run.

6. Results

On the held-out synthetic test split, the classifier reaches **99.81% accuracy**. Training and validation curves are shown in Figure 4; the confusion matrix in Figure 5; per-class metrics in Table 1.

Figure 4. Training/validation accuracy and loss over 100 epochs (reproduction run, seed 42; synthetic data).




Figure 5. Confusion matrix on the synthetic test split (1,037 datasets). Two misclassifications in total.




Table 1. Per-class precision, recall, and F1 on the synthetic test split.

Class	Precision	Recall	F1	Support
KMeans	1.0000	0.9962	0.9981	264
DBSCAN	1.0000	0.9958	0.9979	239
OPTICS	0.9962	1.0000	0.9981	261
GMM	0.9964	1.0000	0.9982	273
Macro avg	0.9981	0.9980	0.9981	1,037
Weighted avg	0.9981	0.9981	0.9981	1,037

Performance is high and near-uniform across classes; the only errors are one KMeans dataset predicted as OPTICS and one DBSCAN dataset predicted as GMM.

7. Discussion

The results show that the 10 hand-designed descriptors carry enough information to separate the four generating families almost perfectly with a small network. In other words, **the feature map ϕ — not the classifier — is where the work is**: once datasets are summarised by these shape and distribution statistics, the four synthetic families become nearly linearly separable, and a 2,500-parameter network suffices. This validates the overall design (describe, then classify) and the choice of descriptors.

It is important to read the 99.81% correctly: it measures how well the model recovers the **generating family** of a synthetic dataset from its descriptors, under the assumption that the family determines the best algorithm. It is not a measurement of clustering quality on real data.

8. Limitations

We are explicit about the boundaries of the result:

1. **Synthetic validation.** All quantitative results are on synthetic data. The companion write-up reports roughly **70%** on real-world datasets, and the original paper states the method was tested mainly on synthetic and only a few real datasets. The synthetic-to-real gap is the principal limitation.
2. **Label assumption.** Each dataset's "best algorithm" label is its generating family, not the empirical winner of running and scoring all four algorithms. This makes the task partly a family-recognition task.
3. **Scope.** The method is restricted to **2-D** data and a **four-algorithm** portfolio.
4. **Pipeline cost.** Feature extraction runs density estimation, contour searches, and several parameter sweeps per dataset, which is computationally heavy at scale.

A candid statement of these limits is part of representing the work honestly: the project is a strong, complete research effort at the secondary-school level, not a production or peer-reviewed system.

9. Conclusion and future work

We presented an end-to-end, learned approach to clustering-algorithm selection: a robust descriptor pipeline (outlier removal, contour-based shape recognition, ten geometric/statistical features) feeding a compact neural classifier, trained on 5,184 balanced synthetic datasets and reaching 99.81% on a held-out synthetic split. Promising directions include: labelling datasets by **empirically** running and scoring all candidate algorithms rather than by generating family; **real-world** evaluation and the synthetic-to-real gap; extending beyond 2-D and to more algorithms; and applying the selector to a concrete downstream task.

Contributions

This was a collaborative team project by Hagen Jacob, Jan Edgar König, and Robert Vetter, developed jointly (largely on a shared notebook, with little formal division of labour, as noted in the original paper). Robert Vetter contributed to the design and implementation of the pipeline and authored the public technical write-up of the project. The external advisors provided scientific guidance but were not authors of the system.

Acknowledgments

We thank our school supervisors **Dr. Marion Moor** and **Johannes Söpke**, and our external scientific advisors **Dr. Wolfgang Felber** and **Christopher Sobel** of **Fraunhofer IIS (Nuremberg)**, who gave regular feedback over the course of the project and hosted a practicum. Fraunhofer's involvement was **advisory only**; the institute did not fund, adopt, or deploy the system.

References

1. J. R. Rice. *The Algorithm Selection Problem*. Advances in Computers, 15:65–118, 1976.
2. S. P. Lloyd. *Least Squares Quantization in PCM*. IEEE Transactions on Information Theory, 28(2):129–137, 1982.
3. J. MacQueen. *Some Methods for Classification and Analysis of Multivariate Observations*. Proc. 5th Berkeley Symposium on Mathematical Statistics and Probability, 1:281–297, 1967.
4. M. Ester, H.-P. Kriegel, J. Sander, X. Xu. *A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise*. Proc. KDD, 226–231, 1996.
5. M. Ankerst, M. M. Breunig, H.-P. Kriegel, J. Sander. *OPTICS: Ordering Points To Identify the Clustering Structure*. Proc. ACM SIGMOD, 49–60, 1999.
6. A. P. Dempster, N. M. Laird, D. B. Rubin. *Maximum Likelihood from Incomplete Data via the EM Algorithm*. Journal of the Royal Statistical Society, Series B, 39(1):1–38, 1977.
7. M. M. Breunig, H.-P. Kriegel, R. T. Ng, J. Sander. *LOF: Identifying Density-Based Local Outliers*. Proc. ACM SIGMOD, 93–104, 2000.
8. P. J. Rousseeuw. *Silhouettes: A Graphical Aid to the Interpretation and Validation of Cluster Analysis*. Journal of Computational and Applied Mathematics, 20:53–65, 1987.
9. E. Parzen. *On Estimation of a Probability Density Function and Mode*. Annals of Mathematical Statistics, 33(3):1065–1076, 1962.
10. M. Rosenblatt. *Remarks on Some Nonparametric Estimates of a Density Function*. Annals of Mathematical Statistics, 27(3):832–837, 1956.
11. F. Pedregosa et al. *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12:2825–2830, 2011.
12. M. Abadi et al. *TensorFlow: A System for Large-Scale Machine Learning*. Proc. USENIX OSDI, 265–283, 2016.

Appendix A — Reproducibility

```
pip install -r requirements.txt
python -m src.data_generation # (optional) regenerate data/data_scaled.xlsx - heavy
python -m src.train           # train (seed 42) -> models/classifier.keras
python -m src.evaluate        # accuracy + per-class report + confusion matrix
```

- **Data:** `data/data_scaled.xlsx` (committed) — the feature table used here.
- **Split:** `train_test_split(test_size=0.2, random_state=42); validation_split=0.2`.
- **Seed:** `tf.keras.utils.set_random_seed(42)` for the reported run.
- **Original notebook:** `notebooks/clustering-algorithm-selection.ipynb` (preserved).
- Regeneration of the dataset uses random states and is computationally heavy, so the regenerated file will not match byte-for-byte; the committed file is the artifact used for the results above.

Appendix B — Feature definitions

See Section 4.4 (Table) for the precise definition of each of the ten features and `../data/DATA.md` for the data schema and class encoding.